

# Online Optimization

Axel Böhm

November 14, 2021

## 1 Introduction

## 2 Strategies

# What is online Learning

Consider the game: In each round  $t = 1, \dots, T$

- ◇ an *adversary* chooses a (secret) number in  $y_t \in [0, 1]$
- ◇ you guess the real number, choosing  $x_t \in [0, 1]$ ;
- ◇ you pay the squared difference  $(x_t - y_t)^2$ .

**Task:** guess a sequence of numbers as precisely as possible.

**Question:** How to measure success?

# Adversary plays i.i.d.

- ◇ Adversary numbers are drawn from a fixed distribution (with mean  $\mu$  and Variance  $\sigma^2$ ).
- ◇ If we knew the distribution, we could pick the mean and pay in expectation  $\sigma^2 T$  (optimal).
- ◇ Benchmark against best possible strategy:

$$\mathbb{E}_Y \left[ \sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 T,$$

or equivalently considering the average

$$\frac{1}{T} \mathbb{E}_Y \left[ \sum_{t=1}^T (x_t - Y)^2 \right] - \sigma^2 .$$

Last quantity should go to zero.

# Minimizing Regret

Let's rewrite a bit more general

$$\mathbb{E} \left[ \sum_{t=1}^T (x_t - Y)^2 \right] - \min_x \mathbb{E} \left[ \sum_{t=1}^T (x - Y)^2 \right] .$$

( $\sigma^2 T$  was just the payoff of the best possible strategy)

Finally:

- ◇ remove the assumption on how the data is generated,
- ◇ consider any arbitrary sequence of  $y_t$  (can remove the expectation).

$$R_T := \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2$$

Called **Regret**, because it measures how much the algorithm regrets for not sticking on all the rounds to the optimal choice in hindsight.

# General loss functions

Online Learning is the study of algorithms to **minimize the regret over a sequence of loss functions** w.r.t. a competitor  $u \in \mathbb{R}^d$ :

$$R_T(u) := \sum_{t=1}^T \ell_t(x_t) - \sum_{t=1}^T \ell_t(u) .$$

This framework allows to

- ◇ reformulate problems in ML and optimization as similar games.
- ◇ analyze situations in which the data are not i.i.d. yet want to guarantee that the algorithm is “learning” something.

For example, online learning can be used to analyze

- ◇ Click prediction problems;
- ◇ spam filter;
- ◇ Convergence to equilibrium of *repeated* games.

It can *also* be used to analyze stochastic algorithms, e.g., SGD.

# Back to the numbers game

Ideally we pick

$$x_T^* := \arg \min_x \sum_{t=1}^T (x - y_t)^2 = \frac{1}{T} \sum_{t=1}^T y_t.$$

- ◇ Don't know the future:  $x_T^*$  is not an option
- ◇ do know the past. in each round: best number **in hindsight**
- ◇ not because we expect the future to be like the past (not true)
- ◇ optimal guess should not change too much between rounds

## Follow-the-Leader (FTL)

Hence, on each round  $t$  our strategy is to guess

$$x_t = x_{t-1}^* = \frac{1}{t-1} \sum_{i=1}^{t-1} y_i.$$

# Follow the leader works for the numbers game

## Theorem

Let  $y_t \in [0, 1]$  for  $t = 1, \dots, T$  an arbitrary sequence of numbers.  
Let the algorithm's output  $x_t = x_{t-1}^* := \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$ . Then,

$$R_T = \sum_{t=1}^T (x_t - y_t)^2 - \min_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2 \leq 4 + 4 \ln T .$$



# Failure of FTL

Let  $V = [-1, 1]$  and consider the sequence of losses  $\ell_t(x) = z_t x$ , where  $z_1 = -0.5$

$$z_t = \begin{cases} 1, & t \text{ even} \\ -1, & t \text{ odd} \end{cases}$$

- ◇ Predictions of FTL will be  $x_t = 1$  for  $t$  even and
- ◇  $x_t = -1$  for  $t$  odd.
- ◇ Cumulative loss of the FTL algorithm will be  $T$ , while fixed solution  $u = 0$  gives 0.
- ◇ regret of FTL is  $T$ .

# Weighted majority algorithm

Consider the **learning from experts** scenario. Experts =  $1, \dots, n$ .  
Decision: "Yes" or "No".

$$f_t(x_t) = \begin{cases} 1 & \text{if wrong} \\ 0 & \text{otherwise} \end{cases}$$

- (i) weights  $w_1(i) = 1$  for all  $i = 1, \dots, n$ , pick  $\alpha > 0$
- (ii) for  $t = 1, \dots, T$ 
  - ❶ compare weights  $\sum_{i \in YES} w_t(i)$  vs.  $\sum_{i \in NO} w_t(i)$
  - ❷ choose Yes or No depending on above comparison
  - ❸ observe feedback
  - ❹ update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if Expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if Expert } i \text{ made a mistake} \end{cases}$$

# Weighted majority algorithm II

Let

- ◇  $M_t$  be the number mistakes we make after  $t$  attempts and
- ◇  $m_t(i)$  the number of mistakes expert  $i$  made (until  $t$ ).

## Theorem

*Then,*

$$M_T \leq 2(1 + \alpha)m_T(i) + 2\frac{\log(n)}{\alpha}$$

*Regret is*

$$M_T - m_T(i^*) = R_T.$$

Approximately: Can bound our mistakes by 2 times number of mistakes of best expert.

# Proof of the Theorem

We always have  $\|w_{t+1}\|_1 \leq \|w_t\|_1$ .

If we made a mistake, then

$$\begin{aligned}\|w_{t+1}\|_1 &\leq \frac{1}{2}\|w_t\|_1 + \frac{1}{2}\|w_t\|_1(1 - \alpha) \\ &= \|w_t\|_1(1 - \alpha/2).\end{aligned}$$

Therefore

$$\|w_{t+1}\|_1 \leq \|w_1\|_1(1 - \alpha/2)^{M_t} = n(1 - \alpha/2)^{M_t}.$$

The weight of  $i$ -th expert can be bounded by

$$w_{t+1}(i) = (1 - \alpha)^{m_t(i)} \leq \|w_{t+1}\|_1$$

Combining the above two yields

$$(1 - \alpha)^{m_t(i)} \leq n(1 - \alpha/2)^{M_t}$$

and

$$m_t(i) \log(1 - \alpha) \leq \log(n) + M_T \log(1 - \alpha/2).$$

## remainder of the proof

Use the fact that for  $x \in (0, \frac{1}{2})$

$$-x - x^2 \leq \log(1 - x) \leq -x$$

to deduce that

$$\begin{aligned} -m_t(i)(\alpha + \alpha^2) &\leq \log(n) - M_T \frac{\alpha}{2} \\ -2m_t(i)(1 + \alpha) &\leq \frac{2}{\alpha} \log(n) - M_T \end{aligned}$$

which yields

$$M_T - \leq 2m_t(i)(1 + \alpha) + \frac{2}{\alpha} \log(n).$$

# Randomized Weighted Majority

Instead of picking the opinion of the (weighted) majority, we only do so with a **probability**.

- (i)  $w_1(i) = 1$  for all  $i = 1, \dots, n$  and  $\alpha \in (0, 1)$
- (ii) for  $t = 1, \dots, T$ 
  - ① compute  $p_t(i) = w_t(i) / \|w_t\|_1$
  - ② choose expert  $i$  with probability  $p_t(i)$
  - ③ observe feedback
  - ④ update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

*Comment:* Randomizing algorithms typically improves the (worst case) analysis.

# Randomized Weighted Majority

Instead of picking the opinion of the (weighted) majority, we only do so with a **probability**.

- (i)  $w_1(i) = 1$  for all  $i = 1, \dots, n$  and  $\alpha \in (0, 1)$
- (ii) for  $t = 1, \dots, T$ 
  - ① compute  $p_t(i) = w_t(i) / \|w_t\|_1$
  - ② choose expert  $i$  with probability  $p_t(i)$
  - ③ observe feedback
  - ④ update weights:

$$w_{t+1}(i) = \begin{cases} w_t(i) & \text{if expert } i \text{ was right} \\ (1 - \alpha)w_t(i) & \text{if expert } i \text{ made a mistake} \end{cases}$$

*Comment:* Randomizing algorithms typically improves the (worst case) analysis.

# Randomized Weighted Majority contd.

As before:

- ◇  $M_t = \#$  of mistakes we make after  $t$  attempts
- ◇  $m_t(i) = \#$  of mistakes expert  $i$  made (until  $t$ )

## Theorem

$$\mathbb{E}[M_T] \leq (1 + \alpha)m_T(i) + \frac{\log(n)}{\alpha}$$

**Improved constants!**



# Multiplicative Weights Algorithm

Before: Loss  $\ell_t$  was 0 or 1.

Now: General loss functions  $\ell_t = (\ell_t(1), \dots, \ell_t(n))$  with  $\ell_t(i) \in [-1, 1]$

(i)  $w_1(i) = 1$  for all  $i = 1, \dots, n$  and  $\alpha \in (0, 1)$

(ii) for  $t = 1, \dots, T$

- ① compute  $p_t(i) = w_t(i) / \|w_t\|_1$
- ② choose expert  $i$  with probability  $p_t(i)$
- ③ observe loss  $\ell_t = (\ell_t(1), \dots, \ell_t(n))$
- ④ update weights:

$$w_{t+1}(i) = (1 - \alpha \ell_t(i)) w_t(i)$$

Note that

$$\langle \mathbf{p}_t, \ell_t \rangle = p_t(1)\ell_t(1) + \dots + p_t(n)\ell_t(n) = \mathbb{E}_i[\ell_t(i)]$$

gives expected loss of round  $t$ .

# Multiplicative Weights Algorithm [contd]

## Theorem

if  $\ell_t(i) \in [-1, 1]$  and  $\alpha < \frac{1}{2}$ , then **MWA** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T |\ell_t(i)| + \frac{\log(n)}{\alpha} \quad \forall i$$

# Hedge Algorithm

- (i)  $w_1(i) = 1$  for all  $i = 1, \dots, n$  and  $\alpha \in (0, 1)$
- (ii) for  $t = 1, \dots, T$ 
  - ① compute  $p_t(i) = w_t(i) / \|w_t\|_1$
  - ② choose expert  $i$  with probability  $p_t(i)$
  - ③ observe loss  $\ell_t = (\ell_t(1), \dots, \ell_t(n))$
  - ④ update weights:

$$w_{t+1}(i) = w_t(i) e^{-\alpha \ell_t(i)}$$

Note:

$$e^{-x} \approx 1 - x$$

# Hedge Algorithm [contd]

## Theorem

If  $\ell_t(i) \in [-1, 1]$  and  $\alpha < \frac{1}{2}$ , then **Hedge Alg.** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T \langle \mathbf{p}_t, \ell_t^2 \rangle + \frac{\log(n)}{\alpha} \quad \forall i.$$

**Further improved constants.**

Observe: Iteration  $t$  is just

$$w_{t+1}(i) = w_t(i) e^{-\alpha \ell_t(i)}$$

$$p_{t+1}(i) = \frac{w_{t+1}(i)}{\|\mathbf{w}_{t+1}\|_1}$$

Online mirror descent! (KL-divergence setting:)

$$h(x) = \sum_i x(i) \log(x(i))$$

# Hedge Algorithm [contd]

## Theorem

If  $\ell_t(i) \in [-1, 1]$  and  $\alpha < \frac{1}{2}$ , then **Hedge Alg.** guarantees

$$\sum_{t=1}^T \langle \mathbf{p}_t, \ell_t \rangle - \sum_{t=1}^T \ell_t(i) \leq \alpha \sum_{t=1}^T \langle \mathbf{p}_t, \ell_t^2 \rangle + \frac{\log(n)}{\alpha} \quad \forall i.$$

**Further improved constants.**

**Observe:** Iteration  $t$  is just

$$w_{t+1}(i) = w_t(i) e^{-\alpha \ell_t(i)}$$

$$p_{t+1}(i) = \frac{w_{t+1}(i)}{\|\mathbf{w}_{t+1}\|_1}$$

Online mirror descent! (KL-divergence setting:)

$$h(x) = \sum_i x(i) \log(x(i))$$

# Online mirror descent

$$x_{t+1} = \arg \min_{x \in K} \{ \langle \nabla f_t(x_t), x \rangle + D_h(x, x_k) \}.$$

Gives bound

$$\sum_{t=1}^T f_t(x_t) - \sum_{t=1}^T f_t(x) \leq \frac{D_h(x, x_0)}{\alpha} + \frac{\alpha T G^2}{2}$$