

# Acceleration of GD via Momentum

Axel Böhm

November 29, 2021

- 1 Optimal methods
- 2 Nesterov momentum
- 3 Heavy ball
- 4 when to use

# Smooth convex functions: less than $\mathcal{O}(\epsilon^{-1})$ steps?

Given  $L$  and  $D = \|x_0 - x^*\|$  we know that **gradient descent**

- ◇ converges with  $\mathcal{O}(1/k)$
- ◇ cannot go faster (“lower bound”)

Maybe GD is not the best possible algorithm?

After all, it is arguably the simplest possible method using the gradient.

# Smooth convex functions: less than $\mathcal{O}(\epsilon^{-1})$ steps?

So let's look at the following classes of methods:

**First-order** method:

- ◇ Access to data only via an **oracle** returning  $f$  and  $\nabla f$  at given points.
- ◇ Clearly, GD is a first order method.

**Q:** What is the **best** first-order method for smooth convex functions.

*best*: smallest upper bound on the number of oracle calls *in the worst case*.

- ◇ Nemirovski and Yudin 1979 proved that

every first-order method needs at least  $\Omega(1/\sqrt{\epsilon})$  iterations  
to find a point  $\bar{x}$  with  $f(\bar{x}) - f^* \leq \epsilon$ .

$\Rightarrow$  no method can be faster than  $\mathcal{O}(1/k^2)$

# Acceleration to $\mathcal{O}(1/\sqrt{\epsilon})$ steps

- ◇ Nesterov 1983 proposed a method that needs only  $\mathcal{O}(1/\sqrt{\epsilon})$  iterations (and is therefore the *best one*).
- ◇ Known as **Nesterov's accelerated gradient** method.
- ◇ By now multiple similar algorithms with same complexity exist.
- ◇ Proofs are generally not really instructive (some are computer assisted).

# Nesterov's accelerated gradient method

---

## Algorithm Nesterov's accelerated gradient method (NAG)

---

```
1: for  $k = 0, 1, \dots$  do
2:    $x_{k+1} = y_k - \frac{1}{L} \nabla f(y_k)$ 
3:    $z_{k+1} = z_k - \frac{k+1}{2L} \nabla f(y_k)$ 
4:    $y_{k+1} = \frac{k+1}{k+3} x_{k+1} + \frac{2}{k+3} z_{k+1}$ 
```

---

- ◇ perform “**smooth step**” from  $y_k$  to  $x_{k+1}$
- ◇ perform **aggressive step** from  $z_k$  to  $z_{k+1}$
- ◇ form **weighted average** of the two  
compensate for the aggressive step by giving less weight

# Nesterov's algorithm as a momentum method

A different way to write the method is via **momentum**

$$\begin{aligned}y_k &= x_k + \beta_k(x_k - x_{k-1}) \\x_{k+1} &= y_k - \frac{1}{L}\nabla f(y_k).\end{aligned}$$

- ◇ differs from GD only in momentum/inertia term  $\beta_k(x_k - x_{k-1})$
- ◇ has to be chosen carefully  $\beta_k = \frac{k-1}{k+2}$
- ◇ coefficient approaches  $\frac{k-1}{k+2} \approx 1 - \frac{3}{k}$

# Nesterov's accelerated gradient method: convergence

Minimum is obtained for  $x^*$ .

## Theorem

Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be convex and  $L$ -smooth, then **NAG** yields

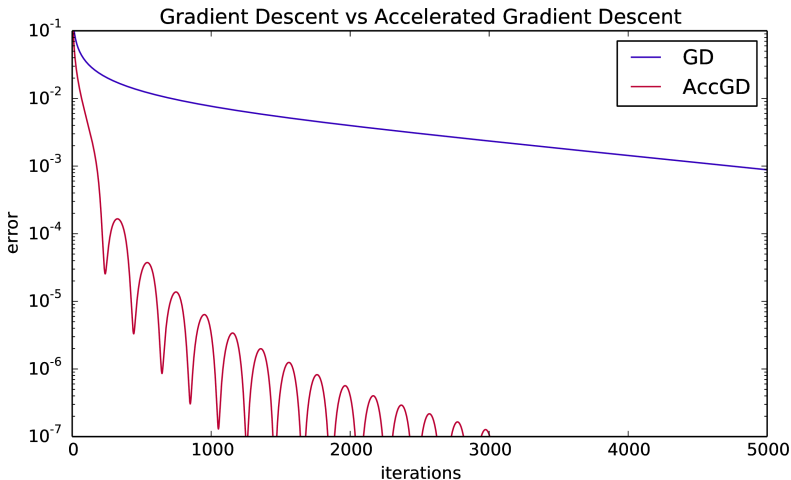
$$f(x_k) - f(x^*) \leq \frac{2L\|x_0 - x^*\|^2}{k(k+1)}$$

Recall that the gradient descent bound was

$$f(x_k) - f(x^*) \leq \frac{L\|x_0 - x^*\|^2}{2k}.$$



# $\mathcal{O}(1/k^2)$ vs $\mathcal{O}(1/k)$ in practice



# Proof idea

Potential function  $\Phi$  that decreases along trajectory (standard technique).  
Out of the blue: Use

$$\Phi(k) := k(k+1)(f(x_k) - f^*) + 2L\|z_k - x^*\|^2.$$

Then show that

$$\Phi(k+1) \leq \Phi(k).$$

Results in

$$\Phi(k+1) \leq \Phi(k) \leq \dots \leq \Phi(0)$$

and therefore

$$k(k+1)(f(x_k) - f^*) \leq 2L\|z_0 - x^*\|^2.$$

# Why momentum?

- ◇ GD has problems with **ravines**, i.e. areas where the surface curves much more steeply in one dimension than in another.
- ◇ Results in zig-zagging.



Figure: no momentum

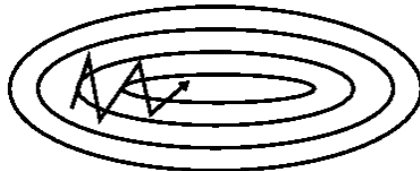
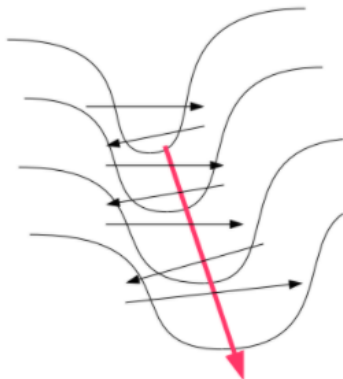
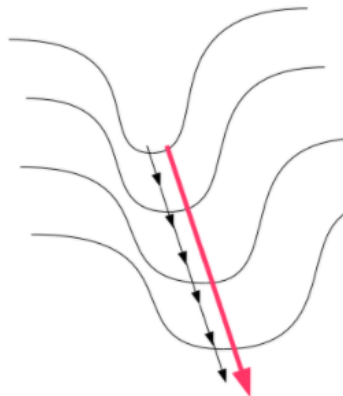


Figure: with momentum

# Momentum and ravines



**SGD bounces back and forth from one side of the valley to the other**



**Using Momentum the zig-zag cancels out, while the direction along the valley is reinforced**

# Momentum in terms of velocity

Consider a ball rolling down a slope. Its **velocity** is

$$v_k = \beta v_{k-1} + \alpha \nabla f(x_k)$$

$$x_{k+1} = x_k - v_k$$

- ◇ a fraction  $\beta$  of the **previous velocity** (friction)
- ◇ plus, steepness of the **slope**

In terms of iterates:

$$x_{k+1} = x_k - v_k$$

$$= x_k - \alpha \nabla f(x_k) - \beta v_{k-1}$$

$$= x_k - \alpha \nabla f(x_k) + \beta(x_k - x_{k-1})$$

# Heavy ball: Polyak 1964

We derived

$$x_{k+1} = x_k - \alpha \nabla f(x_k) + \beta(x_k - x_{k-1}),$$

while Nesterov's method was

$$\begin{aligned} y_k &= x_k + \beta_k(x_k - x_{k-1}) \\ x_{k+1} &= y_k - \frac{1}{L} \nabla f(y_k). \end{aligned}$$

However, **Polyak's** momentum provides no speedup over  $\mathcal{O}(1/k)$  (for smooth convex function).

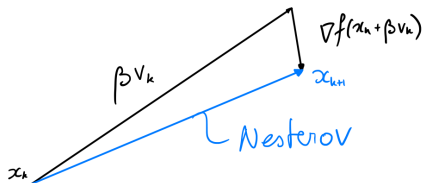
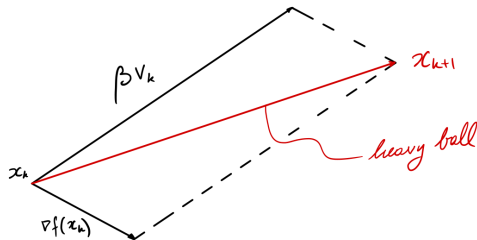
# What's the difference?

- ◇ Both types of momentum seem so similar.
- ◇ Heavy ball does not care if momentum or gradient first.
- ◇ Nesterov momentum applies **inertia first**, then gradient:

$$\begin{aligned}v_k &= \beta v_{k-1} + \alpha \nabla f(x_k + \beta v_{k-1}) \\x_{k+1} &= x_k - v_k.\end{aligned}$$

Provides stabilization if inertia overshoots.

# Nesterov vs. Polyak momentum.





# Momentum for strongly convex functions

For  $L$ -smooth  $\mu$ -strongly convex we know that GD obtains

$$\|x_{k+1} - x^*\|^2 \leq \left(1 - \frac{1}{\kappa}\right) \|x_k - x^*\|^2$$

and

$$f(x_k) - f^* \leq \left(1 - \frac{1}{\kappa}\right)^k \frac{L\|x_0 - x^*\|^2}{2}.$$

Performance depends heavily on the **condition number**  $\kappa := L/\mu$ .  $\Rightarrow$   
Contraction coefficient is  $(1 - 1/\kappa)$ .

Nesterov and Polyak momentum improve this to  $(1 - 1/\sqrt{\kappa})$ .

# Momentum for stochastic methods

**SGD** analysis can be extended to **smooth** functions with rate

$$\mathcal{O}\left(\frac{L}{k} + \frac{\sigma^2}{\sqrt{k}}\right),$$

where  $\sigma^2 := \mathbb{E}[\|\nabla f(x) - g(x)\|^2]$  is the **variance** of the gradient estimator.

This can be improved by Nesterov momentum (and additional tricks) to

$$\mathcal{O}\left(\frac{L}{k^{\textcolor{red}{2}}} + \frac{\sigma^2}{\sqrt{k}}\right).$$

Improvement only in the “**early phase**” before noise takes over.

For worst case rates, only the asymptotic (“late”) phase matters.

# Momentum and nonsmoothness

- ◇ If  $f$  is not differentiable and we have to use subgradients:  
no way to improve the  $\mathcal{O}(1/\sqrt{k})$  rate.
- ◇ Works if objective is **structured**:  $f + g$  (smooth+nonsmooth)

$$\text{FISTA: } \begin{cases} y_k &= x_k + \beta_k(x_k - x_{k-1}) \\ x_{k+1} &= \text{prox}_{\alpha g}(y_k - \alpha \nabla f(y_k)). \end{cases}$$

Accelerates from  $\mathcal{O}(1/k)$  of proximal-gradient method to  $\mathcal{O}(1/k^2)$ .

- ◇ In particular, also works in the **constrained setting**.

# Momentum in the nonconvex world

- ◇ In theory: difficult to show benefit of momentum for nonconvex problems.
  - ▶ some statements under additional smoothness assumptions
- ◇ **Strong empirical evidence** of usefulness.
  - ▶ especially in deep learning.
- ◇ Theory is mostly limited to escaping of saddle points.

Docs &gt; torch.optim &gt; SGD

## SGD

```
CLASS torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0,  

weight_decay=0, nesterov=False) [SOURCE]
```

Implements stochastic gradient descent (optionally with momentum).

---

**input** :  $\gamma$  (lr),  $\theta_0$  (params),  $f(\theta)$  (objective),  $\lambda$  (weight decay),  
 $\mu$  (momentum),  $\tau$  (dampening), *nesterov*

---

**for**  $t = 1$  **to** ... **do**

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$

**if**  $\lambda \neq 0$

$g_t \leftarrow g_t + \lambda \theta_{t-1}$

**if**  $\mu \neq 0$

**if**  $t > 1$

$\mathbf{b}_t \leftarrow \mu \mathbf{b}_{t-1} + (1 - \tau) g_t$

**else**

$\mathbf{b}_t \leftarrow g_t$

**if** *nesterov*

$g_t \leftarrow g_{t-1} + \mu \mathbf{b}_t$

**else**

$g_t \leftarrow \mathbf{b}_t$

$\theta_t \leftarrow \theta_{t-1} - \gamma g_t$

---

**return**  $\theta_t$

---

## Momentum in DL: